

- # Switch as Border Router

Small ISP without FIRT

ITNOG on the WEB 05/2020

nicola modena

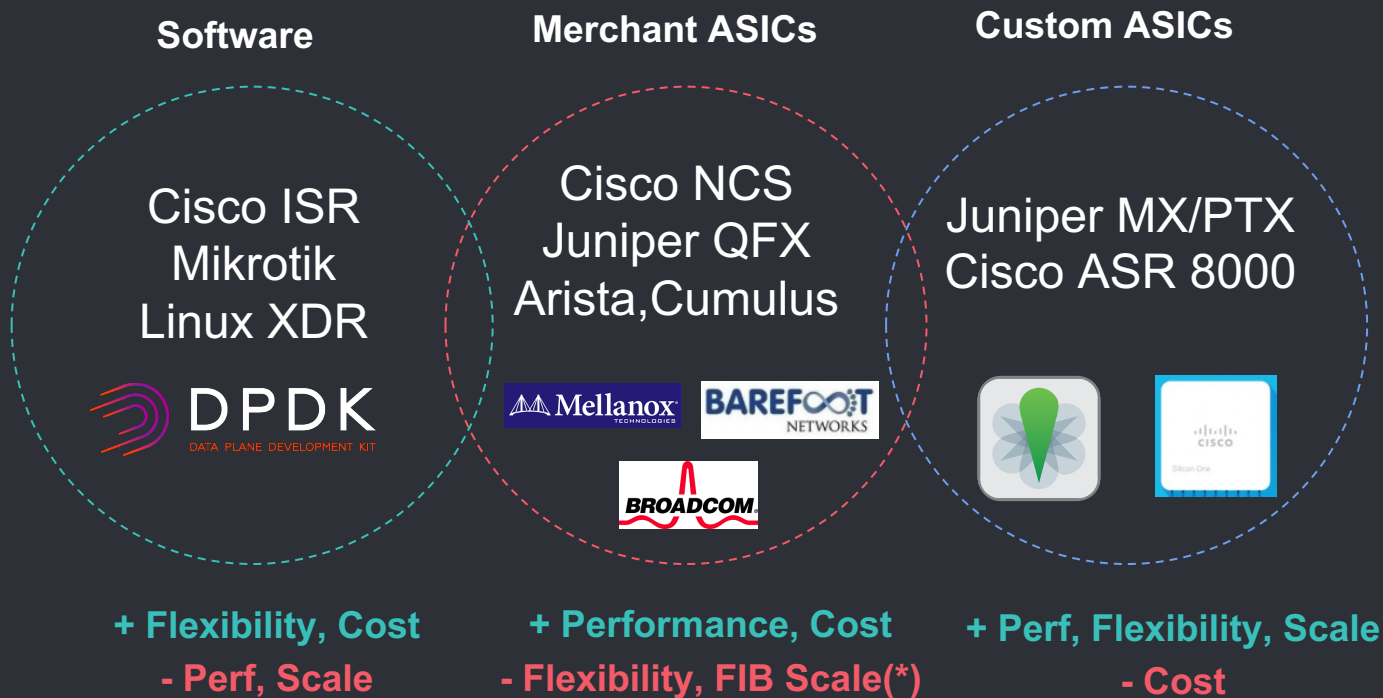
@nmodena



● Agenda

- Motivation
- RIB vs FIB
- IXP and Transit
- Selective routing installation and propagation
- Route Selection for optimized routing
- Self assessment
- Summary

● Software or Hardware based forwarding



- Border router requirement

- Border, Core & Aggregation routers for modern ISP must have:

- High Forwarding capacity
- High speed port density
- High RIB / FIB capacity
- Enhanced Load Balancing
- HW access lists (on the edge)
- Deep buffers (on link speed changes)



RIB vs FIB

Routing Information Base hold all the destination from the best path route selection

One or more entry for each destination are combined and installed in the forwarding table

FIB on ASICs uses fixed size TCAM for efficient lookup.

example: Broadcom Trident III: 128K FIB Entry

Routing Information Base from Best Path Selection

```
SHOW IP ROUTE  
SHOW ROUTE
```

C
J

DRAM



Forwarding Information Base

```
SHOW IP CEF  
SHOW ROUTE FORWARDING-TABLE
```

C

J

TCAM

ISP Backbone without FIRT

use of just default-route to reach external destination (*)

PROS

Lower Resources (FIB)

Faster Convergence

Easier Management

Cheap devices

CONS

sub-optimal routing

No Transit Customers

Possible solution : Route selection & Selective FIB Download

(*) backbone with multiple and distributed exit points

1

A practical example

Border router with peers and transit


IXP Routing

IXP Peering

- eBGP peering with Route Server
- no RS ASN in path
- next-hop unchanged
- BGP update of all members (customizable)

Transit

- FIRT Full Internet Routing Table
- and/or default-route




Our goal is to optimize router resources for this scenario

- Border router with selective FIB install

IXP bgp configuration

```
protocols {
  bgp {
    group IXP {
      neighbor 172.16.0.1 {
        import from-IXP;
        export only-my-AS;
        peer-as 2;
      }
    }
  }
}

policy-options {
  policy-statement from-IXP {
    term accept-all {
      then {
        community add FIB;
      }
    }
  }
  community FIB members 65000:800;
}
```




Tag all the **relevant** prefixes with a community

OPT: filter-out irrelevant destinations reached in other peering points or trough transits

- Border router with selective FIB install

Transit bgp import policy

```
policy-options {  
  policy-statement from-TRANSIT {  
    term default-route {  
      from {  
        route-filter 0.0.0.0/0 exact;  
      }  
      then {  
        community add FIB;  
        next policy;  
      }  
    }  
    term default {  
      then {  
        community add RIB-ONLY;  
      }  
    }  
  }  
  community FIB members 65000:800;  
  community RIB-ONLY members 65000:801;  
}
```




Tag all transit routes with a different community except the default-route.

The FIRT is accepted and may be propagated internally for further usage

- Border router with selective FIB install

Controlled FIB install

```
routing-options {  
  forwarding-table {  
    export selective-FIB-install;  
  }  
}  
  
policy-options {  
  policy-statement selective-FIB-install {  
    term RIB-ONLY {  
      from {  
        protocol bgp;  
        community RIB-ONLY;  
      }  
      then reject;  
    }  
  }  
  community RIB-ONLY members 65000:801;  
}
```



deploy a policy to control RIB to FIB download and reject RIB-ONLY tagged route

HINT: Almost every platform has this feature deployed for out-of-band route reflectors

● Border router with selective FIB install

Checking for the RIB : all entries are present

```
nmodena@MX-05> show bgp summary
```

Table	Tot Paths	Act Paths	Suppressed	History	Damp State	Pending
inet.0						
Peer	AS	InPkt	OutPkt	OutQ	Flaps	Last Up/Dwn
172.16.0.1	2	1294	1304	0	0	9:39:54 Establ
inet.0: 6/6/6/0						
172.16.1.106	1	2073299	9	0	0	53:42 Establ
inet.0: 42793/42793/42793/0						

IXP
TRANSIT


```
nmodena@MX-05> show route
```

```
inet.0: 42822 destinations, 42823 routes (42822 active, 0 holddown, 0 hidden)  
+ = Active Route, - = Last Active, * = Both
```

```
0.0.0.0/0      *[BGP/170] 01:10:09, localpref 100  
                AS path: 1 I, validation-state: unverified  
                > to 172.16.1.106 via et-0/0/4.0  
1.0.0.0/24     *[BGP/170] 01:10:09, localpref 100  
                AS path: 1 49673 24811 13335 I, validation-state: unverified  
                > to 172.16.1.106 via et-0/0/4.0  
1.0.4.0/24     *[BGP/170] 01:10:09, localpref 100  
                AS path: 1 49673 6939 4826 38803 56203 I, validation-state: unverified  
                > to 172.16.1.106 via et-0/0/4.0  
1.0.5.0/24     *[BGP/170] 01:10:09, localpref 100  
                AS path: 1 49673 6939 4826 38803 56203 I, validation-state: unverified  
                > to 172.16.1.106 via et-0/0/4.0
```

FIRT

Border router with selective FIB install



Checking for the FIB : only relevant entries are installed

```
nmodena@MX-05> show route forwarding-table
```

```
Routing table: default.inet
```

```
Internet:
```

```
Enabled protocols: Bridging,
```

Destination	Type	RtRef	Next hop	Type	Index	NhRef	Netif
default	user	0	50::0:b:0:0	ucst	596	4	et-0/0/4.0
default	perm	0		rjct	36	1	
0.0.0.0/32	perm	0		dscd	34	3	
3.0.0.0/8	user	0	172.16.0.3	ucst	599	5	et-0/0/5.0
3.0.1.0/24	user	0	172.16.0.3	ucst	599	5	et-0/0/5.0
3.0.2.0/24	user	0	172.16.0.3	ucst	599	5	et-0/0/5.0
8.0.0.0/8	user	0	172.16.0.8	ucst	598	5	et-0/0/5.0
8.0.1.0/24	user	0	172.16.0.8	ucst	598	5	et-0/0/5.0
8.0.2.0/24	user	0	172.16.0.8	ucst	598	5	et-0/0/5.0
10.0.0.0/8	user	0		dscd	34	3	
10.0.35.0/24	intf	0		rslv	584	1	et-0/0/3.0
10.0.35.0/32	dest	0	10.0.35.0	recv	582	1	et-0/0/3.0
[...]							

Transit

Peers at IXP

Internal Network


● Border router with selective FIB install

Compare RIB and FIB size

```
nmodena@MX-05> show route summary
inet.0: 42822 destinations, 42823 routes (42822 active, 0 holddown, 0 hidden)
    Direct:      5 routes,      5 active
    Local:       4 routes,      4 active
    OSPF:        9 routes,      9 active
    BGP: 42802 routes, 42801 active
    Aggregate:   2 routes,      2 active
    LDP:         1 routes,      1 active

inet.3: 8 destinations, 8 routes (8 active, 0 holddown, 0 hidden)
    LDP:         8 routes,      8 active

nmodena@MX-05> show route forwarding-table summary
Routing table: default.inet
  user:          61 routes
  perm:          5 routes
  intf:          9 routes
  dest:         19 routes
```



Is it really THAT simple ?

yes, but only if you have a SINGLE TRANSIT provider and
all the remaining prefix fits into the FIB
...but this is not true in most cases...

Border router with multiple Transit


Multiple transit are used to get optimal paths like:

- for ASN 20/30/40 use Transit 1
- for ASN 50/60/90 use Transit 2
- load balance the remaining traffic on both links

this requires to

- create policy to identify relevant prefixes
- check if the expected FIB size fits in the TCAM
- mark this new prefixes for FIB download

it is usually a procedure performed manually and if necessary




- Border router and internal reachability

Also internal destination may be considered

and also can be reduced / summarized but depend on

- addressing plan
- summarization strategy
- number of interfaces
- load balancing
- use of MPLS

Hint: some modular devices has also FIB localization capability

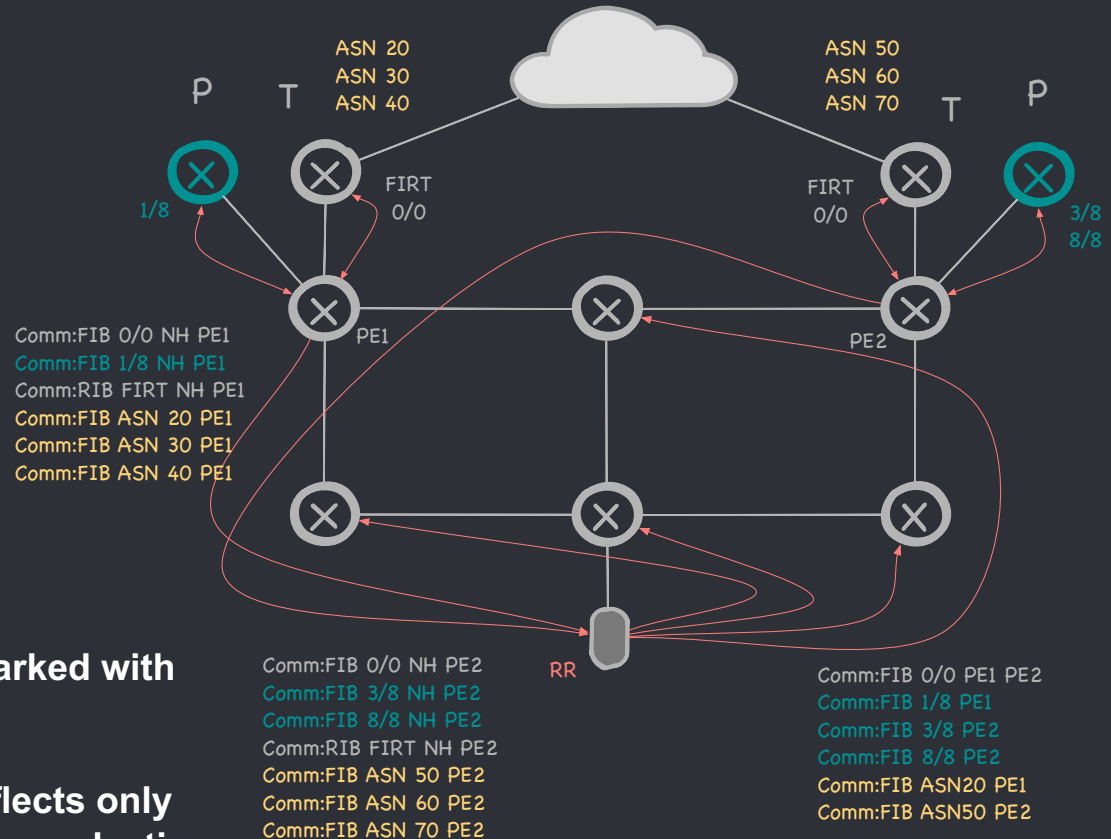


3

Evolution with optimized edge selection

distribution of selected external destinations

Layer-3 solution with optimized routing



We can put all together:

Border routers advertise
all external prefixes to RR marked with
FIB or RIB-ONLY community

RR can hold the FIB and reflects only
relevant prefix to optimize peer selection



What can go wrong ? 😊

- How to identify relevant destinations

Manually : become a regex ninja

Automatic: create an SDN Solution like David Barroso with Paolo Lucente

<https://labs.spotify.com/2016/01/26/sdn-internet-router-part-1/>

<https://blog.ipSPACE.net/2015/01/sdn-router-spotify-on-software-gone-wild.html>

<https://github.com/dbarrosop/sir>

4

Self Assessment

Is this solution good for me ?

● Self Assessment

- how many prefixes I currently have in my rib ?
- how many transit or peering-points I am connected to ?
- what distance (RTT,\$) there is between them ?
- Sub-optimal routing or double transit it's an option ?
- I'm able to identify how many prefix I need to cover 90% of my destination ?
- How many time I spend troubleshooting transit issue ?
- how often I need to optimize transit traffic ?
- I have the knowledge to manage this system ?

-> balance between: TIME - KNOWLEDGE – COSTS - PERFORMANCE

5

Summary

● Switch selection

Select Asic/Vendor based on your needs:

- FIB Size
- Filtering capability
- QOS with Filtering
- Load Balancing
- Buffers
- Netflow
- Telemetry
- Programmability
- CPU / RAM / Control Plane

● Summary - Take Away

- on ASICS resources are fixed/limited
 - We can operate with unsynchronized RIB and FIB
 - Easier on border router with hot-potato and default-route
 - it's possible to deploy on backbone but it's not simple
 - Partial Routing distribution and/or different hasing can lead to traffic loops
 - Troubleshooting requires new skills and updated procedure
-
- you can save a lot of money on high end device

0

THANK YOU

Questions ?

Nicola Modena - CCIE #19119 JNCIE-SP #986
linkedin.com/in/nmodena - @nmodena
nicola@modena.to

Acknowledge:

“Ammiraglio” Tiziano Tofoni friendship and in-depth discussions
Massimo Magnani (Arista) e Roberto Ricci (Cisco) for invaluable inputs
ITNOG community